

Compression Enables Generalization: Wake-Sleep Cycles for Logic Programming with LLM Integration

Alexander Towell

Southern Illinois University Edwardsville

Edwardsville, IL, USA

lex@metafunctor.com

ORCID: 0000-0001-6443-9897

Abstract—Compressing a logic-programming knowledge base, without any labeled positive or negative examples, produces rules that generalize to unseen entities. To isolate this compression signal from LLM memorization, we introduce a test protocol: fact bases with vocabulary the LLM cannot have seen in training, paired with a raw-LLM baseline on the same inputs. On a synthetic crafting domain, direct LLM prompting recovers 0% of held-out generalizations, symbolic compression alone recovers 53%, and a compression pipeline that uses the LLM only as a hypothesis generator for candidate rules recovers $64 \pm 2\%$. Our system, DreamLog, runs wake-sleep cycles inspired by DreamCoder: wake answers queries via SLD resolution with optional LLM rule generation; sleep runs eight compression operations guided by the Minimum Description Length (MDL) principle, each verified against positive and negative queries with atomic rollback. On a canonical family tree, the full pipeline achieves 80% recall with 100% precision. These results are consistent with the Solomonoff induction thesis: shorter programs generalize better.

Index Terms—compression-based learning, logic programming, wake-sleep cycles, knowledge base compression, inductive logic programming, Solomonoff induction, LLM integration

I. INTRODUCTION

Knowledge bases in logic programming accumulate facts over time, yet this accumulation does not constitute learning. A knowledge base with 300 ground facts about a family tree cannot determine whether a newly introduced person is someone’s father, even though the concepts `father`, `mother`, and `grandparent` are implicit in the data. The facts grow, but understanding does not emerge.

This paper presents an empirical test of a classical theoretical claim: *compression is learning*. Solomonoff’s theory of inductive inference [1] and the closely related Minimum Description Length (MDL) principle [2], [3] assert that the shortest program consistent with observations provides the best predictions. We test this claim in a logic programming setting, where “shorter program” means fewer clauses in a knowledge base (KB), and “better predictions” means the ability to derive facts about entities never seen during training.

We introduce DreamLog, a logic programming system that implements *wake-sleep cycles* inspired by DreamCoder [4]. During wake phases, DreamLog answers queries via SLD

resolution, optionally invoking a large language model (LLM) to generate rules for undefined predicates. During sleep phases, eight compression operations transform accumulated facts into general rules. The operations are guided by MDL and verified against a test suite of positive and negative queries, with atomic rollback ensuring behavioral preservation.

The central contribution is an empirical demonstration that KB compression produces rules that generalize to unseen data. We test this on two domains:

- 1) **A novel synthetic domain.** We construct a crafting system with 15 invented materials (“lumite,” “vexal,” “drossite”) that no LLM has encountered in training. Compression discovers concepts such as `master_artisan` and `safe_recipe` purely from structural patterns, achieving 53% recall on unseen entities through symbolic compression alone, and $64 \pm 2\%$ with LLM-assisted operations (a raw LLM baseline achieves 0%).
- 2) **A canonical family tree.** On a 29-person, 4-generation family tree, the full pipeline achieves 80% recall on newly introduced individuals with 100% precision, compressing the KB from 300 to 184 clauses.

The novel domain result is the stronger validation. Because the LLM has never seen the invented terms, it cannot retrieve memorized rules from training data. The symbolic operations discover generalizations through MDL-driven compression of structural regularities in the data, without any labeled training examples or domain-specific inductive bias.

Our contributions are:

- A wake-sleep architecture for logic programming that combines eight symbolic compression operations with LLM-assisted rule discovery, all verified against behavioral test suites.
- Empirical evidence that KB compression enables generalization to unseen entities, including on novel domains with invented vocabulary.
- An ablation study demonstrating complementary contributions: symbolic operations discover entity-level generalizations while the LLM discovers cross-predicate con-

cepts.

- A long-lived accumulation study showing invest-then-compress-then-saturate dynamics consistent with DreamCoder’s convergence.

II. BACKGROUND AND RELATED WORK

A. Compression and Induction

Solomonoff’s theory of inductive inference [1] formalizes the intuition that simpler explanations generalize better. The Kolmogorov complexity of a string is the length of the shortest program that produces it; Solomonoff induction predicts future data by weighting hypotheses inversely by their description length. While Kolmogorov complexity is uncomputable, the MDL principle [2], [3] provides a practical approximation: among models consistent with data, prefer the one with the shortest total description (model plus data given model). In our setting, the “model” is the set of rules in a KB, and the “data given model” is the set of remaining ground facts not derivable from those rules.

Modern empirical work supports the compression-prediction equivalence framing on which we rely. Delétang et al. [5] show that large language models are powerful general-purpose compressors and that the compression viewpoint clarifies scaling laws, tokenization, and in-context learning. Lotfi et al. [6] prove PAC-Bayes generalization bounds via parameter compression that are tight enough to explain why overparameterized neural networks generalize. We do not claim our MDL criterion is tight in the same sense (clause count is a coarse proxy for Kolmogorov complexity), but both lines of work ground the use of compression as a learning signal in measured rather than purely axiomatic terms.

B. DreamCoder and Library Learning

DreamCoder [4], [7] introduced wake-sleep library learning for program synthesis. During wake phases, it solves programming tasks using a library of primitives. During sleep phases, it compresses solved programs to extract reusable abstractions (via anti-unification of lambda terms), which become new library entries. Over iterations, the library converges to a small set of powerful primitives. DreamLog adapts this architecture to logic programming: our “library entries” are Horn clause rules, our “programs” are KB clauses, and our compression operations (particularly Operations C, D, and E) play the role of DreamCoder’s abstraction phase.

Two successors to DreamCoder bear directly on our work. Stitch [8] reformulates the abstraction step as top-down corpus-guided synthesis, achieving 1000–10000× speedup over DreamCoder while preserving library quality. Operations D (predicate invention via skeleton fingerprinting) and E (body pattern extraction) play the role of Stitch for logic programs rather than λ -calculus programs: where Stitch identifies shared λ -term structure across solved tasks, our operations identify shared rule-body structure across the existing KB. LILO [9] extends DreamCoder with LLM-driven documentation and naming of discovered library functions, while keeping the

symbolic compression mechanics. LILO is the closest published spirit to DreamLog: both compress a body of structured artifacts and use the LLM as a complement to symbolic discovery. The two systems differ in what they compress and what the LLM contributes. LILO compresses λ -calculus libraries verified against input/output examples; the LLM names and documents the resulting functions. DreamLog compresses Horn clause knowledge bases verified against deductive closure; the LLM (Operation G) proposes cross-predicate rules the symbolic operations cannot discover. The verification regimes differ correspondingly: LILO checks I/O correctness on held-out tasks, while DreamLog checks query equivalence via the verification suite (Definition 1).

C. Inductive Logic Programming

Inductive Logic Programming (ILP) [10] learns logic programs from labeled positive and negative examples. FOIL [11] uses information gain to grow clause bodies. Progol [12] uses inverse entailment to construct the most specific hypothesis and searches for generalizations. Metagol [13] uses meta-interpretive learning with declarative bias in the form of metarules. ILASP [14] learns Answer Set Programs from partial interpretations. Popper [15] uses hypothesis testing and constraint learning to prune the search space. ∂ ILP [16] and Neural Theorem Provers [17] differentially learn rules from examples. LINC [18] combines LLMs with first-order logic provers for reasoning.

DreamLog differs from ILP in two key respects. First, DreamLog derives its training signal from the KB itself under the closed-world assumption, rather than requiring externally provided labeled examples. All ground facts serve as positive examples, and systematically generated perturbations serve as negative examples. Second, DreamLog’s learning signal is compression, not classification accuracy. Rules are accepted only if they reduce description length while preserving the deductive closure of the KB.

D. Anti-Unification

Anti-unification, introduced by Plotkin [19] and Reynolds [20], computes the least general generalization (lgg) of two terms. It is the dual of Robinson’s unification [21]: where unification finds the most general term subsuming both inputs, anti-unification finds the most specific term that both inputs are instances of. DreamLog uses anti-unification in Operations C and D to identify shared structure across groups of facts and rules.

E. Predicate Invention

Predicate invention [22], [23] creates new predicates not present in the background knowledge. Metagol [13] invents predicates using metarules. Predicate invention is considered one of the most challenging aspects of ILP [24]. DreamLog’s Operation D discovers structurally identical rule sets via skeleton fingerprinting and extracts parameterized predicates using `call/N` dispatch, analogous to lambda abstraction and application.

F. Neural-Symbolic Integration

DeepProbLog [25] integrates neural networks with probabilistic logic programming. Logic Tensor Networks [26] ground logical formulas in real-valued tensors. NeuralLog and related approaches [27], [28] represent relations as learned embeddings. Knowledge graph completion methods (TransE [29], RotatE [30]) learn embedding-based link prediction.

A second cluster combines LLMs with symbolic solvers at inference time. Logic-LM [31] translates natural-language premises into a symbolic form, runs a classical solver, and refines via solver error messages; the LLM acts at parse time and the symbolic engine handles reasoning. LINC [18] follows the same pattern with first-order logic provers. DreamLog takes a different approach: the LLM is used as a *hypothesis generator* during compression (Operation G), proposing rules that the symbolic evaluator verifies before they are added to the KB. This preserves interpretability (every derived fact has a human-readable proof trace through Horn clause resolution) while exploiting LLM priors for cross-predicate patterns, and locates the LLM at a different stage than Logic-LM/LINC: rule *learning* rather than premise translation or input formalization.

G. Never-Ending Learning

NELL [32] continuously learns beliefs from the web, accumulating an ontology over years. DreamLog shares the goal of continuous knowledge accumulation but operates on a different axis: rather than acquiring new facts from external sources, it discovers structure in existing facts through compression. The two approaches are complementary.

III. SYSTEM DESIGN

DreamLog is a logic programming system with S-expression syntax that operates in alternating wake and sleep phases. This section describes the architecture and the eight compression operations that constitute the sleep phase.

A. Wake Phase

During the wake phase, DreamLog accepts assertions and queries. Assertions add ground facts (e.g., `(parent alice bob)`) or user-defined rules to the knowledge base. Queries are answered via SLD resolution [33] with negation-as-failure (NAF) [34], supporting:

- **not/1**: Negation-as-failure with floundering guard.
- **call/N**: Meta-predicate for runtime predicate dispatch.
- **Usage tracking**: Each clause records how often it fires during resolution, providing frequency data for sleep-phase operations.

When a query references an undefined predicate, the system optionally invokes an LLM to propose rules defining that predicate. Generated rules undergo structural validation, optional LLM-as-judge verification, and are added to the KB only if they pass all checks. This mechanism transforms undefined predicates from errors into opportunities for knowledge acquisition.

Algorithm 1 Dream Cycle

Require: Knowledge base K , verification flag v

```

1:  $K_0 \leftarrow \text{copy}(K)$ 
2:  $S \leftarrow \text{BuildVerificationSuite}(K)$  {Pos/Neg queries}
3:  $K \leftarrow \text{OpF}(K)$  {Dead clause pruning (wake data)}
4:  $K \leftarrow \text{OpA}(K)$  {Subsumption elimination}
5:  $K \leftarrow \text{OpB}(K)$  {Redundant fact pruning}
6:  $K \leftarrow \text{OpC}(K, S)$  {Fact generalization}
7:  $S \leftarrow \text{ExtendSuite}(S, K)$  {Add rule-derived queries}
8:  $K \leftarrow \text{OpD}(K, S)$  {Predicate invention}
9:  $K \leftarrow \text{OpE}(K, S)$  {Body pattern extraction}
10:  $K \leftarrow \text{OpG}(K, S)$  {LLM-assisted compression}
11:  $K \leftarrow \text{OpB}(K)$  {Re-prune (post-LLM)}
12:  $K \leftarrow \text{OpH}(K)$  {Lemma caching}
13: if  $v$  and  $\neg \text{Verify}(K, S)$  then
14:    $K \leftarrow K_0$  {Atomic rollback}
15: end if
16: return  $K, |K|/|K_0|$ 

```

B. Sleep Phase: The Dream Cycle

The sleep phase compresses the KB through eight operations, executed in sequence. Algorithm 1 outlines the overall procedure. Operations that modify the KB (C, D, E, G) perform per-candidate verification against the test suite, accepting each candidate only if it preserves behavioral equivalence. A final pipeline-level verification with atomic rollback provides an additional safety net.

Definition 1 (Verification Suite): Given a KB K , the verification suite $S = (S^+, S^-)$ consists of:

- $S^+ = \{t \mid \text{Fact}(t) \in K\}$: all ground facts.
- S^- : synthetic negative queries formed by substituting atoms into existing fact templates at positions where those atoms do not appear in K .

A compression passes verification if every $q \in S^+$ remains derivable and no $q \in S^-$ becomes derivable.

C. Operation A: Subsumption Elimination

Removes clauses subsumed by more general clauses already in the KB. A rule R_1 subsumes R_2 if there exists a substitution θ such that $R_1\theta \sqsubseteq R_2$ and both have the same body length. Additionally, bodyless rules (rules with empty bodies) subsume matching facts. This removes redundancy introduced by prior operations or user input.

D. Operation B: Redundant Fact Pruning

Removes facts that are derivable from the remaining KB (rules plus other facts). Each fact f is tentatively removed; if f remains derivable via SLD resolution, the removal is permanent. Batch removal is attempted first, with one-at-a-time fallback for mutually dependent facts. This is the primary mechanism by which rule discovery translates into fact removal.

E. Operation C: Fact Generalization with Exceptions

This is the central symbolic compression operation. Given a group of $n \geq 3$ facts sharing a functor and arity, the operation:

- 1) **Partitions** facts by argument position: for each position p , groups facts that agree on all arguments except position p .
- 2) **Finds a guard predicate**: a unary predicate g whose extension covers all values appearing at position p in the group (and possibly more).
- 3) **Computes exceptions**: values in the guard’s extension not appearing in the fact group.
- 4) **Applies MDL**: the generalization is accepted only if $1 + |\text{exceptions}| < n$ (one rule plus exception facts is shorter than the original n facts).

Example 2: Given facts `(master_artisan kael)`, `(master_artisan sera)`, ..., `(master_artisan zara)` covering 7 of 9 artisans, and the guard predicate `artisan/1` covering all 9, Operation C produces:

```
(master_artisan X)
:- (artisan X), (not
(exception_master_artisan_artisan
X))
(exception_master_artisan_artisan
thom)
(exception_master_artisan_artisan
fen)
```

This replaces 7 facts with 1 rule + 2 exception facts = 3 clauses, a net reduction of 4.

F. Operation D: Predicate Invention

Discovers structurally identical rule sets across different predicates and extracts a parameterized “invented” predicate. The procedure:

- 1) **Extract skeletons**: For each predicate’s rule set, compute a *skeleton* that abstracts functor names into roles (SELF for recursive calls, `PARAMi` for distinct body functors) while preserving rule structure, arity, and variable connectivity.
- 2) **Group by skeleton**: Predicates with identical skeletons are structurally equivalent.
- 3) **Build parameterized predicate**: Create an invented predicate with an additional `call/N` dispatch parameter, and replace each original predicate with a one-line wrapper delegating to the invented predicate.
- 4) **Apply MDL**: Accept only if $k + m < m \cdot k$, where k is the number of rules per predicate and m is the number of predicates in the group.

This is the logic programming analog of lambda abstraction: the skeleton captures shared computational structure, and `call/N` dispatch plays the role of function application.

G. Operation E: Body Pattern Extraction

Finds common contiguous sub-goal sequences across rule bodies and extracts them as named predicates. Interface variables (those appearing both inside and outside the subse-

quence) become the extracted predicate’s arguments. This discovers shared computational fragments across rules, analogous to common subexpression elimination.

H. Operation F: Dead Clause Pruning

Removes clauses with zero usage after sufficient wake-phase queries. Requires both a minimum query count and that at least 50% of predicates have been exercised. User-provided seed facts are protected from pruning; only clauses added by prior dream cycles (lemmas, LLM-generated rules) are eligible for removal.

I. Operation G: LLM-Assisted Compression

The symbolic operations (A–F) cannot discover rules that cross predicate boundaries (e.g., `father(X,Y) :- parent(X,Y), male(X)`) because they operate within a single functor’s fact group. Operation G invokes an LLM to propose such rules.

The pipeline:

- 1) **Prompt construction**: Sample facts ensuring every predicate is represented (round-robin), include predicate fact counts to guide directionality (specific $\leftarrow$ general).
- 2) **Response parsing**: Parse JSON-formatted rule proposals, handling common LLM formatting errors.
- 3) **Cycle filtering**: Reject rules creating cross-functor cycles in the dependency graph via DFS. Self-recursive rules (depth-limited by the evaluator) are permitted.
- 4) **Helper separation**: Separate helper predicates (new predicates supporting `not/1` patterns) from main rules (deriving existing facts).
- 5) **Evaluation**: Each main rule must derive ≥ 2 existing facts.
- 6) **False-positive check**: Enumerate solutions and reject rules deriving ground terms absent from the KB.
- 7) **Combined verification**: Verify the full set of accepted rules against the verification suite with bounded evaluation to prevent combinatorial explosion.

J. Operation H: Lemma Caching

Adds frequently derived intermediate terms as ground facts for faster resolution. Wake-phase derivation tracking identifies terms computed repeatedly during SLD resolution. Caching these as facts converts multi-step derivations into direct lookups, yielding up to $23.6\times$ query speedup (Section IV).

K. Behavioral Preservation

All operations share a common verification protocol. Before any modification, the system:

- 1) Takes a snapshot of the KB.
- 2) Constructs (or extends) the verification suite.
- 3) Applies the proposed compression.
- 4) Verifies that all positive queries remain derivable and no negative query becomes derivable.
- 5) Rolls back to the snapshot if verification fails.

This ensures that the compressed KB is *behaviorally equivalent* to the original on all tested queries. Experiment EX02

confirmed zero semantic drift across 200 held-out queries over multiple dream cycles.

IV. EXPERIMENTS

We evaluate DreamLog on four axes: generalization to unseen entities (the core claim), comparison against a modern ILP baseline (Popper), long-lived knowledge accumulation, and ablation of the compression pipeline. All experiments use Anthropic Claude Haiku 4.5 for Operation G, with total LLM cost under \$0.10 across all experiments.

A. Experimental Setup

All experiments run on DreamLog’s Python implementation with SLD resolution, NAF, and `call/N`. The LLM provider for Operation G is Anthropic Claude Haiku 4.5 (\$0.25/M input tokens) via API. Each dream cycle invokes the LLM at most twice (once for rule proposal in Operation G, once for predicate naming). Verification suites are constructed automatically from KB state. All compression operations use default parameters: minimum group size 3, shared structure threshold 0.1, maximum 200 prompt facts.

B. EX25b: Generalization on a Novel Domain

1) *Motivation*: The canonical test for compression-driven generalization (a family tree) is problematic: the LLM has likely seen `father(X, Y) :- parent(X, Y), male(X)` thousands of times during training. We cannot distinguish compression-driven discovery from training data recall. This experiment uses a synthetic domain with invented terms the LLM has never encountered.

2) *Domain*: An alchemical crafting workshop with 15 materials having invented names (lumite, vexal, drossite, pyreth, quarl, noctite, aerine, sylphex, thalline, morven, zephore, ethylis, fenrik, glacine, bramith) and properties (phase, material class, hazard status). The domain includes 16 recipes combining materials into products, 9 artisans with skills, and 5 skill requirements. Total: 112 base facts and 61 derived facts across 6 predicates (`hazardous_recipe`, `safe_recipe`, `same_phase_recipe`, `metallic_alloy`, `can_craft`, `master_artisan`).

3) *Protocol*: Four conditions: (1) *no dream*: baseline KB with no compression; (2) *symbolic*: dream cycle with Operations A–F only (no LLM); (3) *full pipeline*: dream cycle with all operations including LLM-assisted Operation G; (4) *raw LLM*: prompt Haiku directly with all KB facts and new-entity queries (no compression pipeline). After dreaming, 6 new materials, 8 new recipes, and 4 new artisans are introduced with base facts only. 33 checks (19 positive, 14 negative) test whether derived predicates are correctly inferred for the unseen entities. LLM-dependent conditions are run 5 times; we report mean $\pm$ standard deviation.

TABLE I
GENERALIZATION ON THE NOVEL CRAFTING DOMAIN (EX25B). 33 CHECKS (19 POSITIVE, 14 NEGATIVE) ON UNSEEN ENTITIES. LLM CONDITIONS: MEAN $\pm$ STD OVER 5 RUNS.

Condition	Rules	Acc	Prec	Recall
No dream	0	42%	—	0%
Raw LLM	0	42%	—	0%
Symbolic	5	52%	59%	53%
Full	20	58 $\pm$ 1%	64 $\pm$ 1%	64 $\pm$ 2%

4) *Results*: Table I shows the results.

The *raw LLM* baseline achieves 0% recall: Haiku answers “NO” to every query about invented terms, confirming that direct LLM prompting contributes nothing on novel domains. The compression pipeline is genuinely necessary.

Symbolic compression alone discovers 5 rules and achieves 53% recall. The key discoveries are:

- `master_artisan(X) :- artisan(X), not(exception...)`: generalizes from the pattern that 7 of 9 artisans have 2+ skills.
- `safe_recipe(X) :- product(X), not(exception...)`: generalizes from the pattern that most recipes use non-hazardous materials.

These are discovered by Operation C through MDL-driven fact generalization. No domain knowledge is used; the system finds that “most artisans are master artisans” and “most recipes are safe” purely from the statistical structure of the data.

The full pipeline (mean of 5 runs) achieves 64 $\pm$ 2% recall with 20 discovered rules. Operation G adds cross-predicate rules including `hazardous_recipe` (via material hazard properties) and concepts like `volatile_material` and `metallic_material`. The low variance ($\pm$ 2%) confirms that results are stable across LLM generations.

Seven false positives occur from `safe_recipe` and `same_phase_recipe` generalizations applying to exception cases (e.g., a liquid-solid mix classified as same-phase). This illustrates the inherent risk of inductive generalization.

The undiscovered predicate is `can_craft(Artisan, Product)`, which requires a 3-hop join (`artisan  $\rightarrow$  skill  $\rightarrow$  requires_skill  $\rightarrow$  recipe_type`). This exceeds the current operations’ reach.

C. EX25: Generalization on a Canonical Domain

1) *Domain*: A 4-generation family tree with 29 people, 96 base facts (`person`, `male/female`, `parent`), and 204 derived facts across 7 predicates (`father`, `mother`, `grandparent`, `grandfather`, `grandmother`, `great_grandparent`, `ancestor`).

2) *Protocol*: Same conditions as EX25b, including the raw LLM baseline. After dreaming, 6 new people are introduced with base facts only, and 14 checks (10 positive, 4 negative) test derived relationships.

3) *Results*: Table II shows the results.

The full pipeline discovers 7 rules: `father  $\leftarrow$  parent + male`, `mother  $\leftarrow$  parent + female`, `grandparent  $\leftarrow$`

TABLE II

GENERALIZATION ON THE CANONICAL FAMILY TREE (EX25). THE FULL PIPELINE ACHIEVES 80% RECALL WITH 100% PRECISION.

Condition	Rules	Acc	Prec	Recall	KB Size
No dream	0	29%	—	0%	300
Symbolic	1	29%	—	0%	—
Full	7	86%	100%	80%	184

TABLE III

CROSS-DOMAIN COMPARISON OF GENERALIZATION RECALL. RAW LLM SCORES 0% ON BOTH DOMAINS.

Domain	No Dream	Raw LLM	Symbolic	Full
Crafting (novel)	0%	0%	53%	64 ± 2%
Family (canonical)	0%	0%	0%	80%

parent ◦ parent, grandfather ← grandparent + male, grandmother ← grandparent + female, and variants. The KB compresses from 300 to 184 clauses (ratio 0.613). Precision is 100%: no false positives.

The sole unrecovered predicate is ancestor, which requires a recursive rule (ancestor(X, Y) :- parent(X, Y) and ancestor(X, Z) :- parent(X, Y), ancestor(Y, Z)). The current LLM prompt does not propose recursive rules, and symbolic operations cannot discover them from flat fact patterns.

Symbolic compression discovers only one entity-specific rule (e.g., “albert is an ancestor of everyone”) rather than the general concept. This highlights the critical role of the LLM: symbolic operations find patterns *within* a predicate, while the LLM discovers relationships *across* predicates.

D. Comparison: Novel vs. Canonical Domain

Table III compares the two domains directly.

The novel domain provides the stronger evidence for compression-driven learning. The symbolic operations achieve 53% recall on a domain with invented terms, demonstrating that MDL-guided generalization works without any prior knowledge. The family domain achieves higher overall recall (80%) but relies entirely on the LLM (Operation G), making it difficult to disentangle compression-driven discovery from LLM training data memorization.

E. EX26: Popper ILP Baseline

1) *Motivation*: DreamLog and Popper [15] occupy different points in the rule-learning design space. Popper is supervised (it requires positive and negative examples plus mode declarations) and uses a constraint-driven search; DreamLog is unsupervised (it consumes an unannotated fact base) and uses a compression-driven search. A naive head-to-head would advantage one system or the other depending on how labels are supplied. We construct an apples-to-apples comparison that gives each system the natural form of its required inputs and evaluates both on the same held-out facts.

TABLE IV

EX26 POPPER BASELINE. APPLES-TO-APPLES COMPARISON ON THE SAME EX25 FAMILY AND EX25B CRAFTING DOMAINS, THREE SEEDS PER CELL. DREAMLOG CONDITIONS MATCH EX25/EX25B NUMBERS WITHIN ROUNDING; POPPER IS THE NEW DATA POINT.

Domain	System	Recall	Prec	Acc	Rules
Family	No dream	0%	—	29%	0
	Symbolic	0%	—	29%	1
	Full	80%	100%	86%	7
	Popper	80%	67–80%	57–71%	6
Crafting	No dream	0%	—	42%	0
	Symbolic	53%	59%	52%	5
	Full	63%	63%	58%	21–24
	Popper	42%	80%	61%	4

2) *Protocol*: For each target predicate, the same EX25/EX25b fact base is converted to Popper input format: `bk.pl` contains the base facts plus the other targets’ positives; `exs.pl` contains the held-in derived facts of the target predicate as E^+ and a closed-world random sample (with $|E^-| = |E^+|$, drawn with the same RNG seed as the data split) as E^- ; `bias.pl` contains hand-written mode declarations from the natural schema. Modes are committed before any results are observed and not tuned after the fact. Per-cell timeout is 60 seconds; failures (timeout, parse error, crash) are recorded as zero rules. We exclude the recursive ancestor target: Popper’s Python-level timeout does not interrupt its clingo C-extension on this configuration, and EX25 also does not recover ancestor, so both systems get a pass on the recursive case. Three random seeds (42, 43, 44) per (system, domain) cell, yielding 24 cells total (4 systems × 2 domains × 3 seeds). Total wall time was approximately four minutes; LLM cost \$0.03 across the six full-pipeline cells.

3) *Results*: Table IV reports recall, precision, accuracy, and rule count for each cell, averaged over the three seeds.

On the family domain, Popper matches DreamLog’s full-pipeline recall (both 80%) but trails on precision by 20–33 percentage points. The rules tell the story: Popper discovers `father( $X, Y$ ) :- parent( $X, Y$ )` and `mother( $X, Y$ ) :- parent( $X, Y$ )` (and analogous over-general grandfather/grandmother rules) because random closed-world negatives do not include “parent but female” near-misses that would force the gender constraint. The two correct rules Popper discovers are `grandparent` and `great_grandparent`, where no auxiliary predicate is needed.

On the crafting domain, the result inverts. Popper’s recall (42%) is *below* DreamLog’s symbolic-only baseline (53%) by 11 percentage points, and well below the full pipeline (63%). Both Popper and DreamLog-symbolic are pure symbolic learners with no LLM, so the gap argues that DreamLog’s MDL-guided fact generalization (Operation C) and skeleton-based abstraction (Operations D, E) have an inductive bias that random-negative ILP does not naturally express. Popper compensates with higher precision (80% vs. DreamLog full’s

63%), producing four conservative rules rather than 20+ exploratory ones.

4) *Interpretation:* We do not read these numbers as “DreamLog beats Popper” or vice versa. The two systems answer different signals. The findings characterize the methodological tradeoff cleanly. On family, the closed-world random-negative protocol is too easy for Popper because the universe of false “father” tuples is dominated by clearly-impossible pairs (e.g., a child as parent of their grandparent); a curated set of near-miss negatives would close the gap. On crafting, hand-curating equivalent near-misses is not natural because the invented vocabulary does not come with a pre-existing intuition about what should be excluded. DreamLog’s contribution is operational rather than competitive: when only an unannotated fact base is available, LLM-supplied semantic priors (Operation G) and MDL-guided structural generalization (Operations C, D, E) substitute for the curated supervision Popper relies on.

F. EX25c: Holdout Sweep and the Open-World Mode

1) *Motivation:* The new-entity tests (EX25, EX25b) introduce entities with no prior facts at all. A complementary protocol, common in inductive logic programming, holds out a fraction of derived facts and asks whether the system can recover them after compressing the remainder. This sweep characterizes the data density threshold for compression-driven generalization on the family tree domain.

2) *Closed-World versus Open-World Verification:* Operation G’s verification pipeline includes a false-positive check: any rule that derives a ground term absent from the KB is rejected. This is correct for the standard compression setting (a complete KB should not be augmented with novel facts) but pathological for holdout evaluation, where the goal is precisely to recover absent facts.

We introduce an *open-world mode* that disables this check. In this mode, Op G accepts rules that cover at least 2 existing facts and pass the verification suite, regardless of whether the rule also derives facts not in the KB. The closed-world default preserves the strong safety guarantee for normal use; open-world mode is enabled explicitly during holdout evaluation.

3) *Protocol:* For each ratio $r \in \{0.05, 0.10, 0.20, 0.30, 0.50\}$, we randomly hold out r of each predicate’s derived facts (stratified by predicate). The remaining facts plus all base facts form the training KB. We dream on the training KB (full pipeline, open-world mode) and measure recovery: the fraction of held-out facts derivable from the compressed KB. Three random seeds per ratio.

4) *Results:* Table V shows the recovery curve.

Three observations:

Open-world recovers all held-out facts at low ratios.

At 5–20% holdout, recovery is 100%: the discovered rules (father $\leftarrow$ parent + male, etc.) generalize perfectly to the held-out facts. The KB is compressed from 290 to 118 clauses (ratio 0.40), an even stronger compression than the no-holdout setting in EX25 (0.61), because the smaller KB has less verification overhead.

TABLE V

HOLDOUT SWEEP ON FAMILY TREE (EX25C). 3 SEEDS PER RATIO. CLOSED-WORLD RECOVERY IS 0% AT EVERY RATIO (THE FP CHECK REJECTS RULES THAT DERIVE ANY HELD-OUT FACT). OPEN-WORLD RECOVERY FOLLOWS A CLEAR CURVE.

Holdout	Closed-world	Open-world	Compression
5%	0%	100%	0.405
10%	0%	100%	0.399
20%	0%	100%	0.415
30%	0%	$85 \pm 21\%$	0.507
50%	0%	$72 \pm 20\%$	0.660

The data density threshold appears at 30%. Recovery drops to $85 \pm 21\%$ at 30% holdout and $72 \pm 20\%$ at 50% holdout. The high variance reflects which facts get held out: some random splits leave enough facts per predicate for Op G to propose a rule that derives at least 2 existing facts (Op G’s acceptance gate); others do not. On the family tree, the cross-predicate rules that drive recovery (father $\leftarrow$ parent + male, etc.) come from Op G rather than Op C, so the relevant constraint is Op G’s minimum-coverage gate, not Op C’s `min_group_size`.

No false positives in either mode. Even with the FP check disabled, the rules discovered by Op G do not derive incorrect facts. This is because the rules themselves (father=parent+male) are sound; the closed-world FP check was rejecting them not for being wrong but for deriving facts that were intentionally absent.

The closed-world default remains the right choice for production use (strong safety guarantee, zero semantic drift). Open-world mode is an explicit opt-in for evaluation against ILP-style benchmarks where the closed-world assumption is not part of the protocol.

G. EX24: Long-Lived Knowledge Accumulation

1) *Motivation:* Real systems accumulate knowledge incrementally over sessions. We test whether DreamLog exhibits meaningful learning dynamics over an extended interaction.

2) *Protocol:* A simulated research lab scenario: “Dr. Chen” builds institutional knowledge across 25 sessions spanning two simulated years. Seven domains: people, publications, teaching, grants, infrastructure, collaborations, and service. Sessions vary from 5 to 30 assertions. Session 10 includes corrections (role changes). Session 17 includes user-provided rules with NAF helper predicates. Dream cycles occur every 5 sessions.

3) *Results:* The final KB contains 451 clauses (422 facts, 29 rules) from 443 total assertions. Table VI shows the dream cycle progression.

Three phases are visible:

1) **Investment** (Dream 2): The LLM discovers 17 cross-domain rules including `grant_funded_work`, `cross_institution_collaboration` (a 4-body rule with NAF), and `mentorship_chain`. The KB *expands* from 179 to 195 (ratio 1.09), investing in structural rules.

TABLE VI
DREAM CYCLE PROGRESSION IN THE 25-SESSION RESEARCH LAB (EX24).

Cycle	Session	Before	After	Key Discoveries
1	S05	94	91	2 symbolic rules
2	S10	179	195	17 LLM cross-domain rules
3	S15	278	275	2 symbolic rules
4	S20	—	—	Saturation (minor pruning)
5	S25	—	—	No change

TABLE VII
LEAVE-ONE-OUT ABLATION ON A 119-CLAUSE KB (EX08).
PERCENTAGES ARE NOT ADDITIVE; EACH ROW SHOWS THE COMPRESSION
LOST WHEN THAT OPERATION ALONE IS DISABLED.

Operation	Contribution	Role
C: Generalization	85% (11 clauses)	Primary compressor
B: Fact pruning	15% (2 clauses)	Post-rule cleanup
D: Predicate invention	15% (2 clauses)	Structural abstraction
E: Body extraction	−15% (+2 clauses)	Adds structural clarity
A, F, H	0%	Situational

- 2) **Compression** (Dream 3): Symbolic operations capitalize on accumulated rules, achieving net compression.
- 3) **Saturation** (Dreams 4–5): No new patterns remain. The KB is fully compressed.

This invest-then-compress-then-saturate dynamic matches DreamCoder’s convergence behavior [4]. The 29 discovered rules come from three sources: 4 symbolic (Operation C), 17 LLM (Operation G), and 8 user-provided. All 13 correctness checks pass, including transitive citation chains, NAF-based active membership, and negative examples. Total LLM cost: \$0.023.

H. Ablation Analysis

Table VII summarizes the ablation from Experiment EX08 on a 119-clause KB with diverse pattern types.

Operation C is the workhorse, contributing 85% of compression on this KB. Operation E *increases* clause count (it adds a clause for the extracted predicate) but improves structural clarity, which may benefit future compression cycles. Operations A and F contribute on KBs with subsumed or dead clauses, and Operation H contributes to query performance rather than KB size.

The generalization experiments (Tables I and II) provide a complementary ablation: symbolic operations (C primarily) discover entity-level generalizations (“most artisans are masters”), while Operation G discovers cross-predicate concepts (“father = male parent”). Neither alone achieves the full pipeline’s performance.

I. Additional Results

1) *Convergence (EX01)*: The dream cycle converges in exactly 2 iterations across all tested KBs: cycle 1 compresses, cycle 2 confirms no further opportunities exist. The cycle is a monotone operator on KB size that converges quickly in practice.

2) *Semantic Drift (EX02)*: Zero drift across 200 held-out queries over all dream cycles. The verification suite with positive and negative queries prevents any behavioral change. This is a strong safety guarantee for production use.

3) *DreamCoder Benchmark (EX03)*: On list-processing programs (member, append, reverse, length, map, filter), the system discovers `forall/2`: a parameterized predicate equivalent to DreamCoder’s “every” combinator. This validates the architectural analogy between the two systems.

4) *Scaling (EX09)*: Compression ratio is stable at 0.88 across KBs from 69 to 1019 clauses (10 to 200 entities). Throughput degrades from 75 clauses/s to 12 clauses/s due to quadratic verification cost.

5) *Noise Tolerance (EX11)*: Robust to 0–20% noise (incorrect facts left as isolated facts, not generalized). At 30%+ noise, Operation C generalizes erroneous patterns that share sufficient structural regularity.

6) *Query Speedup (EX10)*: After dreaming, query time drops from 25.0ms to 1.1ms (23.6× speedup). Operation H caches 15 frequently derived lemma facts, converting multi-step recursive resolutions into direct lookups.

V. DISCUSSION

A. What Symbolic vs. LLM Operations Contribute

The experiments reveal a clean division of labor. Symbolic operations (particularly Operation C) discover generalizations *within* a predicate: “most artisans are master artisans,” “most recipes are safe.” These are statistical regularities visible from the distribution of ground facts. The LLM (Operation G) discovers relationships *across* predicates: “father = parent who is male,” “hazardous recipe = recipe using a hazardous material.” These require understanding how distinct predicates relate semantically.

On the novel crafting domain, symbolic compression achieves 53% recall while the full pipeline reaches 64%. On the canonical family tree, symbolic compression achieves 0% recall (it finds only entity-specific patterns) while the full pipeline achieves 80%. Critically, a raw LLM baseline (prompting Haiku directly with all facts) achieves 0% on both domains, confirming that the compression pipeline provides the generalization capability, not the LLM alone. This asymmetry reflects the nature of the domains: the crafting domain has within-predicate regularities (most artisans are masters) while the family domain’s generalization opportunities are entirely cross-predicate.

B. The Novel Domain Isolates Compression from LLM Memorization

The crafting domain result provides stronger evidence for compression-driven generalization than the family tree. When the LLM proposes `father(X, Y) :- parent(X, Y), male(X)` for a family KB, it may be recalling this rule from training data rather than discovering it from structural patterns. The invented terms in the crafting domain (lumite, vexal, drossite) eliminate this confound. That symbolic compression alone achieves 53% recall on invented terms, while the raw

LLM achieves 0%, demonstrates genuine MDL-driven concept discovery independent of any LLM prior knowledge.

The family tree remains useful as a benchmark because its higher recall (80%) demonstrates the ceiling achievable when the LLM can draw on domain familiarity. The two domains together bracket the system’s capability: novel domains test pure compression, canonical domains test the LLM integration.

C. Limitations

1) *Multi-hop rules*: The system cannot discover rules requiring 3+ hop joins. `can_craft(Artisan, Product)` requires chaining `artisan → skill → requires_skill → recipe_type`, which exceeds the body length of rules currently discoverable by either symbolic or LLM operations.

2) *Recursive rules*: The ancestor predicate `(ancestor(X, Z) :- parent(X, Y), ancestor(Y, Z))` is not discovered. Symbolic operations work on flat fact patterns; the LLM prompt does not currently encourage recursive proposals. Supporting recursive rule discovery is an important direction for future work.

3) *Minimum data threshold*: The holdout sweep (EX25c) shows recovery degrades sharply once more than 30% of facts per predicate are held out. Op G accepts a proposed rule only if it derives at least 2 existing facts, so rules become harder to discover as the remaining fact count per predicate shrinks. Op C’s generalizations also weaken (its default `min_group_size` is 3), but on the family tree the recovery curve is driven by Op G’s cross-predicate rules. Compression-driven learning needs sufficient data density per predicate to activate; the threshold is roughly 7–10 facts per predicate for the family tree domain.

4) *Scaling*: Verification cost is quadratic in KB size due to exhaustive query checking. KBs beyond 1000 clauses would benefit from sampled verification or cached resolution.

5) *LLM dependence for cross-predicate rules*: The family tree ablation shows that symbolic operations alone achieve 0% generalization recall when all opportunities are cross-predicate. The system relies on the LLM for this critical capability, introducing dependence on LLM quality and availability.

D. Relation to Inductive Logic Programming

DreamLog occupies a different point in the rule learning design space than classical ILP systems. We briefly contrast with three representative approaches, clarifying what DreamLog does and does not claim.

a) *Popper [15]*: treats rule learning as a constraint satisfaction problem. Given explicit positive and negative examples, Popper searches for a hypothesis that covers all positives and excludes all negatives, using failures on counter-examples as constraints to prune the search space. The learning signal is classification accuracy on labeled data. DreamLog’s signal is instead compression: rules are accepted when they reduce description length while preserving deductive closure. Section IV-E (EX26) reports the apples-to-apples comparison: Popper supplied with the same KB plus closed-world-sampled negatives matches DreamLog’s full-pipeline recall on

the family domain (both 80%) but trails by 20–33 precision points (over-general rules missing gender constraints); on the novel crafting domain Popper’s recall (42%) falls below DreamLog’s symbolic-only baseline (53%). The novel-domain result is the more discriminating one: both systems are pure symbolic learners there, and the gap argues that DreamLog’s MDL-guided structural generalization has an inductive bias random-negative ILP does not naturally express. Popper’s constraint-driven search and DreamLog’s compression-driven search remain complementary signals that could in principle be combined.

b) *Metagol [13]*: uses meta-interpretive learning with user-provided metarules as declarative bias. The metarules constrain the hypothesis space and make predicate invention tractable. DreamLog’s Operation D invents predicates via skeleton fingerprinting without metarules, relying instead on structural equivalence across rule sets already present in the KB. Metagol can solve problems DreamLog cannot (e.g., learning recursive predicates from a few examples when the metarule “closure” is provided), and DreamLog can solve problems Metagol cannot (e.g., discovering that seven of nine artisans share a pattern without needing anyone to write a metarule). Again, the signals are different rather than competitive.

c) *LINC [18]*: uses an LLM as a translator: natural language premises are converted to first-order logic, then a classical theorem prover decides entailment. LINC’s LLM operates at parse time, not inference time. DreamLog uses an LLM differently: Operation G treats the LLM as a *hypothesis generator* for compression candidates, which are then verified by a symbolic evaluator. Both approaches preserve symbolic soundness by keeping the reasoning engine classical, but they apply the LLM at different stages (input formalization vs. rule proposal). LINC does not learn rules; it reasons over given ones. The two systems address different problems.

In summary, we do not claim that DreamLog is a better ILP system than Popper or Metagol. We claim that compression-driven learning from unannotated fact bases is a distinct capability, that it produces rules which generalize to unseen entities, and that combining symbolic compression with LLM-assisted rule proposal yields complementary benefits. EX26 provides a direct head-to-head against Popper; extending the comparison to canonical ILP benchmarks (Michalski trains [10], Bongard problems, NELL [32]) remains future work and would require accepting explicit positive/negative example sets alongside fact bases.

E. Relation to Solomonoff Induction

Our results are consistent with the Solomonoff thesis. The compressed KBs generalize better than the uncompressed originals. However, our MDL criterion (clause count) is a coarse proxy for Kolmogorov complexity, and two domains do not establish a predictive relationship between compression ratio and generalization recall. A more faithful implementation would account for clause complexity (number of variables,

body length) rather than treating all clauses as unit cost. We leave this refinement to future work.

VI. CONCLUSION

We have presented DreamLog, a logic programming system that discovers generalizable concepts through knowledge base compression. The central finding is that compression enables generalization: a KB compressed from 300 to 184 clauses derives facts about entities it has never seen, including on novel domains with invented vocabulary. Symbolic compression alone achieves 53% recall on such domains, while a raw LLM baseline achieves 0%, validating MDL-driven concept discovery independent of any LLM prior knowledge. The full pipeline, combining symbolic and LLM-assisted operations, achieves up to 80% recall with 100% precision on canonical domains.

The system exhibits meaningful learning dynamics over extended interactions: an invest-then-compress-then-saturate pattern consistent with DreamCoder’s convergence behavior, suggesting a general property of compression-based learning systems.

Several directions remain open. Direct comparison with ILP systems (ILASP, Metagol) on shared benchmarks would position DreamLog precisely in the landscape of rule learners. Support for recursive rule discovery would address the ancestor predicate gap. Transfer learning experiments (training on one domain, compressing a combined KB with a second) would test whether compression discovers cross-domain abstractions. Scaling beyond 1000 clauses requires sampled verification. Finally, a richer MDL metric incorporating clause complexity (variable count, body length, nesting depth) would bring the system closer to Solomonoff’s theoretical ideal.

The code, all experiment scripts, and the experiment registry with full provenance are available at <https://github.com/quaelius/dreamlog>.

REFERENCES

- [1] R. J. Solomonoff, “A formal theory of inductive inference. part i,” *Information and control*, vol. 7, no. 1, pp. 1–22, 1964.
- [2] J. Rissanen, “Modeling by shortest data description,” *Automatica*, vol. 14, no. 5, pp. 465–471, 1978.
- [3] P. D. Grünwald, *The Minimum Description Length Principle*. MIT Press, 2007.
- [4] K. Ellis, C. Wong, M. Nye, M. Sablé-Meyer, L. Morales, L. Hewitt, L. Cary, A. Solar-Lezama, and J. B. Tenenbaum, “Dreamcoder: Bootstrapping inductive program synthesis with wake-sleep library learning,” in *Proceedings of the 42nd ACM SIGPLAN International Conference on Programming Language Design and Implementation (PLDI)*. ACM, 2021, pp. 835–850.
- [5] G. Delétang, A. Ruoss, P.-A. Duquenne, E. Catt, T. Genewein, C. Matern, J. Grau-Moya, L. K. Wenliang, M. Aitchison, L. Orseau, M. Hutter, and J. Veness, “Language modeling is compression,” in *International Conference on Learning Representations (ICLR)*, 2024.
- [6] S. Lotfi, M. Finzi, S. Kapoor, A. Potapczynski, M. Goldblum, and A. G. Wilson, “PAC-Bayes compression bounds so tight that they can explain generalization,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.
- [7] K. Ellis, L. Wong, M. Nye, M. Sablé-Meyer, L. Cary, L. Anaya Pozo, L. Hewitt, A. Solar-Lezama, and J. B. Tenenbaum, “DreamCoder: growing generalizable, interpretable knowledge with wake-sleep Bayesian program learning,” *Philosophical Transactions of the Royal Society A: Mathematical, Physical and Engineering Sciences*, vol. 381, no. 2251, p. 20220050, 2023.
- [8] M. Bowers, T. X. Olausson, L. Wong, G. Grand, J. B. Tenenbaum, K. Ellis, and A. Solar-Lezama, “Top-down synthesis for library learning,” *Proceedings of the ACM on Programming Languages*, vol. 7, no. POPL, pp. 1182–1213, 2023.
- [9] G. Grand, L. Wong, M. Bowers, T. X. Olausson, M. Liu, J. B. Tenenbaum, and J. Andreas, “LILO: Learning interpretable libraries by compressing and documenting code,” in *International Conference on Learning Representations (ICLR)*, 2024.
- [10] S. Muggleton and L. De Raedt, “Inductive logic programming: Theory and methods,” *The Journal of Logic Programming*, vol. 19, pp. 629–679, 1994.
- [11] J. R. Quinlan, “Learning logical definitions from relations,” *Machine learning*, vol. 5, no. 3, pp. 239–266, 1990.
- [12] S. Muggleton, “Inverse entailment and prolog,” *New generation computing*, vol. 13, no. 3, pp. 245–286, 1995.
- [13] S. H. Muggleton, D. Lin, N. Pahlavi, and A. Tamaddoni-Nezhad, “Meta-interpretive learning: application to grammatical inference,” *Machine Learning*, vol. 94, no. 1, pp. 25–49, 2014.
- [14] M. Law, A. Russo, and K. Broda, “ILASP: Learning answer set programs from examples,” in *Proceedings of the 24th International Conference on Logic Programming*, 2014.
- [15] A. Cropper and R. Morel, “Learning programs by learning from failures,” *Machine Learning*, vol. 110, no. 4, pp. 801–856, 2021.
- [16] R. Evans and E. Grefenstette, “Learning explanatory rules from noisy data,” *Journal of Artificial Intelligence Research*, vol. 61, pp. 1–64, 2018.
- [17] T. Rocktäschel and S. Riedel, “End-to-end differentiable proving,” in *Advances in Neural Information Processing Systems*, 2017, pp. 3788–3800.
- [18] T. X. Olausson, A. Gu, B. Lipkin, C. E. Zhang, A. Solar-Lezama, J. B. Tenenbaum, and R. Levy, “LINC: A neurosymbolic approach for logical reasoning by combining language models with first-order logic provers,” in *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, 2023, pp. 5153–5176.
- [19] G. D. Plotkin, “A note on inductive generalization,” *Machine Intelligence*, vol. 5, pp. 153–163, 1970.
- [20] J. C. Reynolds, “Transformational systems and the algebraic structure of atomic formulas,” *Machine Intelligence*, vol. 5, pp. 135–151, 1970.
- [21] J. A. Robinson, “A machine-oriented logic based on the resolution principle,” *Journal of the ACM*, vol. 12, no. 1, pp. 23–41, 1965.
- [22] I. Stahl, “Predicate invention in inductive logic programming,” *Advances in Inductive Logic Programming*, pp. 34–47, 1995.
- [23] S. Muggleton and W. Buntine, “Machine invention of first-order predicates by inverting resolution,” in *Proceedings of the 5th International Conference on Machine Learning*, 1988, pp. 339–352.
- [24] A. Cropper, S. Dumancic, R. Evans, and S. H. Muggleton, “Inductive logic programming at 30: a new introduction,” *Journal of Artificial Intelligence Research*, vol. 74, pp. 765–850, 2022.
- [25] R. Manhaeve, S. Dumancic, A. Kimmig, T. Demeester, and L. De Raedt, “DeepProbLog: Neural probabilistic logic programming,” in *Advances in Neural Information Processing Systems*, 2018, pp. 3749–3759.
- [26] L. Serafini and A. d’Avila Garcez, “Logic tensor networks: Deep learning and logical reasoning from data and knowledge,” in *Neural Symbolic Learning and Reasoning (NeSy)*, 2016.
- [27] F. Yang, Z. Yang, and W. W. Cohen, “Differentiable learning of logical rules for knowledge base reasoning,” in *Advances in Neural Information Processing Systems*, 2017, pp. 2319–2328.
- [28] P. W. Battaglia, J. B. Hamrick, V. Bapst, A. Sanchez-Gonzalez, V. Zambaldi, M. Malinowski, A. Tacchetti, D. Raposo, A. Santoro, R. Faulkner et al., “Relational inductive biases, deep learning, and graph networks,” *arXiv preprint arXiv:1806.01261*, 2018.
- [29] A. Bordes, N. Usunier, A. Garcia-Duran, J. Weston, and O. Yakhnenko, “Translating embeddings for modeling multi-relational data,” in *Advances in Neural Information Processing Systems*, 2013, pp. 2787–2795.
- [30] Z. Sun, Z.-H. Deng, J.-Y. Nie, and J. Tang, “RotatE: Knowledge graph embedding by relational rotation in complex space,” in *International Conference on Learning Representations*, 2019.
- [31] L. Pan, A. Albalak, X. Wang, and W. Y. Wang, “Logic-LM: Empowering large language models with symbolic solvers for faithful logical reasoning,” in *Findings of the Association for Computational Linguistics: EMNLP 2023*, 2023.
- [32] T. Mitchell, W. Cohen, E. Hruschka, P. Talukdar, B. Yang, J. Betteridge, A. Carlson, B. Dalvi, M. Gardner, B. Kisiel et al., “Never-ending learning,” *Communications of the ACM*, vol. 61, no. 5, pp. 103–115, 2018.

- [33] J. W. Lloyd, *Foundations of Logic Programming*, 2nd ed. Springer, 1987.
- [34] K. L. Clark, "Negation as failure," *Logic and Data Bases*, pp. 293–322, 1978.